Module 3 Lab Exercise: Machine Learning Workflow and Types of Learning

Learning Objectives

By the end of this lab, you will be able to:

- · Distinguish between supervised, unsupervised, and reinforcement learning
- Understand the complete machine learning workflow
- · Build and evaluate your first classification model
- · Work with different types of data (numerical, categorical, text, images)
- Apply the end-to-end ML process: data → model → evaluation → insights

Prerequisites

- · Completed Module 2 (familiar with Python libraries and Jupyter/Colab)
- · Understanding of basic data operations and visualization
- · Access to your GitHub repository for saving work

Part 1: Understanding Types of Machine Learning

Machine learning can be categorized into three main types. Let's explore each with practical examples.

1. Supervised Learning

Definition: Learning from labeled examples to make predictions on new, unseen data.

Examples:

- Classification: Predicting categories (spam/not spam, disease/healthy)
- Regression: Predicting continuous values (house prices, temperature)

Key Characteristic: We have both input features (X) and correct answers (y) during training.

2. Unsupervised Learning

Definition: Finding hidden patterns in data without labeled examples.

Examples:

- Clustering: Grouping similar customers for marketing
- · Dimensionality Reduction: Simplifying complex data while keeping important information

Key Characteristic: We only have input features (X), no correct answers during training.

3. Reinforcement Learning

Definition: Learning through trial and error by receiving rewards or penalties.

Examples:

- · Game playing (chess, Go)
- Autonomous vehicles
- · Recommendation systems that learn from user feedback

Key Characteristic: Agent learns by interacting with an environment and receiving feedback.

For this course, we'll focus primarily on supervised learning, with some unsupervised learning in later modules.

Part 2: Setting Up Our Machine Learning Environment

Let's start by importing our libraries and loading a dataset that will help us understand the ML workflow.

```
# Import essential libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```



```
import seaborn as sns
from sklearn.datasets import load wine, make classification
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
# Set style for better-looking plots
plt.style.use('default')
sns.set_palette("husl")
print("
All libraries imported successfully!")
print("
    Ready to start our machine learning journey!")
   All libraries imported successfully!
   Ready to start our machine learning journey!
```

Part 3: Loading and Exploring Our Dataset

We'll use the Wine dataset - a classic dataset for classification. It contains chemical analysis of wines from three different cultivars (types) grown in Italy.

```
# Load the Wine dataset
wine data = load wine()
# Convert to DataFrame for easier handling
df = pd.DataFrame(wine_data.data, columns=wine_data.feature_names)
df['wine_class'] = wine_data.target
df['wine_class_name'] = [wine_data.target_names[i] for i in wine_data.target]
print("Dataset Information:")
print(f"Shape: {df.shape}")
print(f"Features: {len(wine_data.feature_names)}")
print(f"Classes: {wine_data.target_names}")
print(f"\nFirst 5 rows:")
print(df.head())
Dataset Information:
Shape: (178, 15)
Features: 13
Classes: ['class_0' 'class_1' 'class_2']
First 5 rows:
   alcohol malic_acid ash alcalinity_of_ash magnesium total_phenols \
   14.23
                 1.71 2.43
                                         15.6
                                                   127.0
                 1.78 2.14
1
    13.20
                                         11.2
                                                   100.0
                                                                   2.65
    13.16
                 2.36 2.67
                                          18.6
                                                   101.0
                                                                   2.80
    14.37
                 1.95 2.50
                                          16.8
                                                                   3.85
                                                   118.0
                                                                   2.80
    13.24
                 2.59 2.87
                                         21.0
   flavanoids nonflavanoid_phenols proanthocyanins color_intensity hue
0
        3.06
                              0.28
                                              2.29
                                                              5.64 1.04
1
        2.76
                              0.26
                                              1.28
                                                               4.38 1.05
2
        3.24
                              0.30
                                              2.81
                                                              5.68 1.03
3
        3.49
                              0.24
                                              2.18
                                                               7.80 0.86
4
        2.69
                              0.39
                                              1.82
                                                               4.32 1.04
   od280/od315_of_diluted_wines proline wine_class_wine_class_name
                          3.92 1065.0
                                                         class 0
                          3.40 1050.0
1
                                                 0
                                                           class 0
2
                          3.17
                                1185.0
                                                 0
                                                          class_0
                          3.45 1480.0
                                                          class_0
4
                          2.93
                                 735.0
                                                 0
                                                           class_0
```

```
# Explore the dataset structure
print("Dataset Overview:")
print("=" * 50)
print(f"Total samples: {len(df)}")
print(f"Features (input variables): {len(df.columns) - 2}") # -2 for target columns
print(f"Target classes: {df['wine_class_name'].unique()}")
print(f"\nClass distribution:")
print(df['wine_class_name'].value_counts())
```

```
# Check for missing values
print(f"\nMissing values: {df.isnull().sum().sum()}")
print("
No missing values - this is a clean dataset!")
Dataset Overview:
_____
Total samples: 178
Features (input variables): 13
Target classes: [np.str_('class_0') np.str_('class_1') np.str_('class_2')]
Class distribution:
wine_class_name
class_1
class_0
          59
class_2
          48
Name: count, dtype: int64
Missing values: 0
✓ No missing values - this is a clean dataset!
```

Part 4: Exploratory Data Analysis (EDA)

Before building models, we need to understand our data. This is a crucial step in the ML workflow.

```
# Visualize class distribution
plt.figure(figsize=(12, 4))
# Subplot 1: Class distribution
plt.subplot(1, 2, 1)
class_counts = df['wine_class_name'].value_counts()
plt.bar(class_counts.index, class_counts.values, color=['red', 'green', 'blue'])
plt.title('Distribution of Wine Classes')
plt.xlabel('Wine Class')
plt.ylabel('Number of Samples')
plt.xticks(rotation=45)
# Subplot 2: Feature correlation heatmap (first 6 features for clarity)
plt.subplot(1, 2, 2)
correlation_matrix = df.iloc[:, :6].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0, fmt='.2f')
plt.title('Feature Correlations (First 6 Features)')
plt.tight_layout()
plt.show()
print("[] EDA helps us understand:")
print("- Class balance (are all classes equally represented?)")
print("- Feature relationships (which features are correlated?)")
print("- Data quality (any outliers or issues?)")
                       Distribution of Wine Classes
                                                                                          Feature Correlations (First 6 Features)
                                                                                                                                            1.0
   70
                                                                              alcohol -
                                                                                                0.09
                                                                                                        0.21
                                                                                                                -0.31
                                                                                                                        0.27
                                                                                                                                0.29
                                                                                                                                            0.8
   60
Number of Samples
                                                                           malic acid -
                                                                                       0.09
                                                                                                1.00
                                                                                                        0.16
                                                                                                                0.29
                                                                                                                        -0.05
                                                                                                                                -0.34
                                                                                                                                            0.6
   50
                                                                                                0.16
                                                                                 ash
                                                                                       0.21
                                                                                                        1.00
                                                                                                                0.44
                                                                                                                         0.29
                                                                                                                                0.13
                                                                                                                                            0.4
   40
                                                                                                        0.44
                                                                                       -0.31
                                                                                                0.29
                                                                                                                        -0.08
                                                                                                                                -0.32
                                                                      alcalinity_of_ash -
   30
                                                                                                                                            0.2
   20
                                                                                                                -0.08
                                                                          magnesium -
                                                                                       0.27
                                                                                                -0.05
                                                                                                        0.29
                                                                                                                                 0.21
                                                                                                                                            0.0
   10
                                                                                                                         0.21
                                                                                                                                            -0.2
                                                                         total_phenols -
                                                                                                -0.34
                                                                                                        0.13
                                                                                                                -0.32
                                  1855
              1855)
                                                                                         alcohol
                                                                                                         ash
                                                                                                                 alcalinity of ash
                                                                                                                                  total_phenols
                                                                                                 malic_acid
                                                                                                                          magnesium
                                Wine Class
EDA helps us understand:
- Class balance (are all classes equally represented?)
- Feature relationships (which features are correlated?)
- Data quality (any outliers or issues?)
```

Part 5: The Complete Machine Learning Workflow

Now let's implement the standard ML workflow step by step:

The 6-Step ML Workflow:

- 1. Data Preparation: Clean and prepare the data
- 2. Feature Selection: Choose relevant input variables
- 3. Data Splitting: Separate training and testing data
- 4. Model Training: Teach the algorithm using training data
- 5. Model Evaluation: Test performance on unseen data
- 6. Model Interpretation: Understand what the model learned

Let's implement each step!

```
# Step 1: Data Preparation
print("Step 1: Data Preparation")
print("=" * 30)
# Select features (X) and target (y)
# For simplicity, let's use the first 4 features
feature_names = ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash']
X = df[feature names]
y = df['wine_class']
print(f"Selected features: {feature_names}")
print(f"Feature matrix shape: {X.shape}")
print(f"Target vector shape: {y.shape}")
# Display first few rows
print("\nFirst 5 samples:")
print(X.head())
Step 1: Data Preparation
Selected features: ['alcohol', 'malic acid', 'ash', 'alcalinity of ash']
Feature matrix shape: (178, 4)
Target vector shape: (178,)
First 5 samples:
   alcohol malic_acid ash alcalinity_of_ash
   14.23
             1.71 2.43
                                          15.6
    13.20
                 1.78 2.14
                                          11.2
    13.16
                 2.36 2.67
                                          18.6
3
    14.37
                 1.95 2.50
                                          16.8
    13.24
                 2.59 2.87
                                          21.0
```

```
# Step 2: Data Splitting
print("Step 2: Data Splitting")
print("=" * 30)
# Split data into training (80%) and testing (20%) sets
X_train, X_test, y_train, y_test = train_test_split(
   Χ, ν,
   test_size=0.2,
                      # 20% for testing
   random_state=42,  # For reproducible results
    stratify=y
                      # Maintain class proportions
print(f"Training set: {X_train.shape[0]} samples")
print(f"Testing set: {X_test.shape[0]} samples")
print(f"Training classes: {np.bincount(y_train)}")
print(f"Testing classes: {np.bincount(y_test)}")
print("\n@ Why split data?")
print("- Training set: Teach the model")
print("- Testing set: Evaluate performance on unseen data")
print("- This prevents overfitting (memorizing vs. learning)")
Step 2: Data Splitting
_____
Training set: 142 samples
Testing set: 36 samples
Training classes: [47 57 38]
Testing classes: [12 14 10]
```

```
Why split data?Training set: Teach the modelTesting set: Evaluate performance on unseen dataThis prevents overfitting (memorizing vs. learning)
```

```
# Step 3: Model Training
print("Step 3: Model Training")
print("=" * 30)
# Create and train two different models
    'Logistic Regression': LogisticRegression(random_state=42),
    'Decision Tree': DecisionTreeClassifier(random_state=42, max_depth=3)
trained_models = {}
for name, model in models.items():
    print(f"\nTraining {name}...")
   # Train the model
   model.fit(X_train, y_train)
   trained_models[name] = model
   print("\n What happened during training?")
print("- Models learned patterns from training data")
print("- They found relationships between features and wine classes")
print("- Now they can make predictions on new data!")
Step 3: Model Training
_____
Training Logistic Regression...
✓ Logistic Regression training completed!
Training Decision Tree...
Decision Tree training completed!
What happened during training?
- Models learned patterns from training data
- They found relationships between features and wine classes
- Now they can make predictions on new data!
```

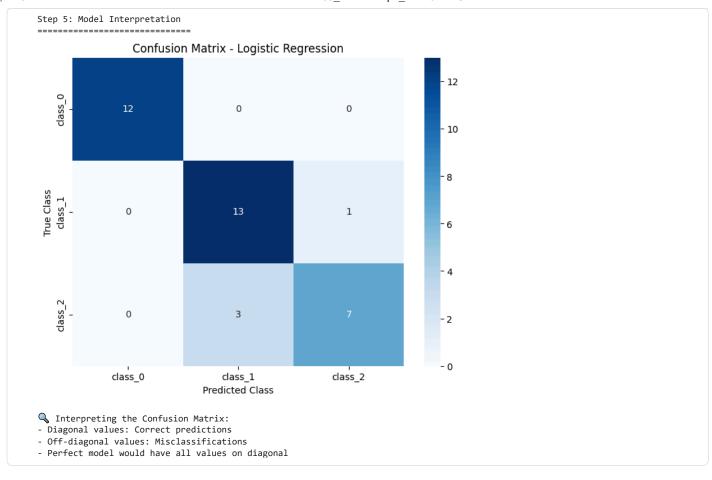
```
# Step 4: Model Evaluation
print("Step 4: Model Evaluation")
print("=" * 30)
results = {}
for name, model in trained_models.items():
    # Make predictions
    y_pred = model.predict(X_test)
    # Calculate accuracy
    accuracy = accuracy_score(y_test, y_pred)
    results[name] = accuracy
    print(f"\n{name} Results:")
    print(f"Accuracy: {accuracy:.3f} ({accuracy*100:.1f}%)")
    # Detailed classification report
    print("\nDetailed Performance:")
    print(classification_report(y_test, y_pred, target_names=wine_data.target_names))
# Compare models
print("\n Model Comparison:")
for name, accuracy in results.items():
    print(f"{name}: {accuracy:.3f}")
best_model = max(results, key=results.get)
print(f"\n
    Best performing model: {best_model}")
Step 4: Model Evaluation
```

```
Logistic Regression Results:
Accuracy: 0.889 (88.9%)
Detailed Performance:
             precision
                         recall f1-score
                                     1.00
    class_0
                  1.00
                           1.00
                                                12
    class_1
                  0.81
                           0.93
                                     0.87
                                                14
    class_2
                  0.88
                           0.70
                                     0.78
                                                10
                                     0.89
   accuracy
                                                36
  macro avg
                  0.90
                           0.88
                                     0.88
                                                36
weighted avg
                  0.89
                           0.89
                                     0.89
                                                36
Decision Tree Results:
Accuracy: 0.833 (83.3%)
Detailed Performance:
                         recall f1-score
             precision
                                           support
    class_0
                  0.86
                           1.00
                                     0.92
                                                12
    class_1
                  0.91
                           0.71
                                     0.80
                                                14
    class_2
                  0.73
                                     0.76
                           0.80
                                                10
   accuracy
                                     0.83
                                                36
  macro avg
                  0.83
                           0.84
                                     0.83
                                                36
weighted avg
                  0.84
                           0.83
                                     0.83
                                                36

■ Model Comparison:

Logistic Regression: 0.889
Decision Tree: 0.833
```

```
# Step 5: Model Interpretation
print("Step 5: Model Interpretation")
print("=" * 30)
# Visualize confusion matrix for the best model
best_model_obj = trained_models[best_model]
y_pred_best = best_model_obj.predict(X_test)
plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_pred_best)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            \verb|xticklabels=wine_data.target_names|,\\
            yticklabels=wine_data.target_names)
plt.title(f'Confusion Matrix - {best_model}')
plt.xlabel('Predicted Class')
plt.ylabel('True Class')
plt.show()
print(f"\n \ Interpreting the Confusion Matrix:")
print("- Diagonal values: Correct predictions")
print("- Off-diagonal values: Misclassifications")
print("- Perfect model would have all values on diagonal")
```



Part 6: Understanding Different Data Types in ML

Machine learning works with various types of data. Let's explore the main categories:

```
# Understanding Different Data Types in ML
print("Understanding Data Types in Machine Learning")
print("=" * 45)
# Create examples of different data types
data examples = {
    'Numerical (Continuous)': [23.5, 45.2, 67.8, 12.1, 89.3],
    'Numerical (Discrete)': [1, 5, 3, 8, 2],
    'Categorical (Nominal)': ['Red', 'Blue', 'Green', 'Red', 'Blue'],
    'Categorical (Ordinal)': ['Low', 'Medium', 'High', 'Medium', 'Low'],
    'Text': ['Hello world', 'Machine learning', 'Data science', 'Python programming', 'AI revolution'],
    'Boolean': [True, False, True, True, False]
for data_type, examples in data_examples.items():
   print(f"\n{data_type}:")
   print(f" Examples: {examples}")
   print(f" Use case: ", end="")
   if 'Continuous' in data_type:
       print("Regression problems (predicting prices, temperatures)")
    elif 'Discrete' in data_type:
       print("Counting problems (number of items, ratings)")
   elif 'Nominal' in data_type:
       print("Classification without order (colors, categories)")
    elif 'Ordinal' in data_type:
        print("Classification with order (ratings, sizes)")
   elif 'Text' in data_type:
       print("Natural language processing (sentiment analysis, translation)")
   elif 'Boolean' in data_type:
        print("Binary classification (yes/no, spam/not spam)")
print("\n♥ Key Insight: Different data types require different preprocessing and algorithms!")
```

```
Understanding Data Types in Machine Learning
Numerical (Continuous):
 Examples: [23.5, 45.2, 67.8, 12.1, 89.3]
 Use case: Regression problems (predicting prices, temperatures)
Numerical (Discrete):
 Examples: [1, 5, 3, 8, 2]
 Use case: Counting problems (number of items, ratings)
Categorical (Nominal):
 Examples: ['Red', 'Blue', 'Green', 'Red', 'Blue']
 Use case: Classification without order (colors, categories)
Categorical (Ordinal):
 Examples: ['Low', 'Medium', 'High', 'Medium', 'Low']
 Use case: Classification with order (ratings, sizes)
Text:
  Examples: ['Hello world', 'Machine learning', 'Data science', 'Python programming', 'AI revolution']
 Use case: Natural language processing (sentiment analysis, translation)
Boolean:
 Examples: [True, False, True, True, False]
 Use case: Binary classification (yes/no, spam/not spam)
Rey Insight: Different data types require different preprocessing and algorithms!
```

Part 7: Hands-On Practice - Build Your Own Model

Now it's your turn! Complete the following tasks to reinforce your learning.

```
# Task 1: Try different features
print("Task 1: Experiment with Different Features")
print("=" * 40)
# Your task: Select 3 different features and build a model
# Available features:
print("Available features:")
for i, feature in enumerate(wine_data.feature_names):
    print(f"{i+1:2d}. {feature}")
# TODO: Replace these with your chosen features
your_features = ['alcohol', 'color_intensity', 'proline'] # Modify this list
# Build model with your features
X_your = df[your_features]
X_train_your, X_test_your, y_train_your, y_test_your = train_test_split(
    X_your, y, test_size=0.2, random_state=42, stratify=y
# Train a logistic regression model
your_model = LogisticRegression(random_state=42)
your_model.fit(X_train_your, y_train_your)
# Evaluate
y_pred_your = your_model.predict(X_test_your)
your_accuracy = accuracy_score(y_test_your, y_pred_your)
print(f"\nYour model features: {your_features}")
print(f"Your model accuracy: {your_accuracy:.3f} ({your_accuracy*100:.1f}%)")
# Compare with original model
print(f"Original model accuracy: {results['Logistic Regression']:.3f}")
if your_accuracy > results['Logistic Regression']:
   print(" See Great job! Your feature selection improved the model!")
    print("P Try different features to see if you can improve performance!")
Task 1: Experiment with Different Features
Available features:
1. alcohol
2. malic_acid
3. ash
4. alcalinity_of_ash
```

```
5. magnesium
6. total_phenols
7. flavanoids
8. nonflavanoid_phenols
9. proanthocyanins
10. color_intensity
11. hue
12. od280/od315_of_diluted_wines
13. proline

Your model features: ['alcohol', 'color_intensity', 'proline']
Your model accuracy: 0.833 (83.3%)
Original model accuracy: 0.889

Try different features to see if you can improve performance!
```

Part 8: Assessment - Understanding ML Concepts

Answer the following questions to demonstrate your understanding:

```
# Assessment Task 1: Identify the ML type
print("Assessment Task 1: Identify Machine Learning Types")
print("=" * 50)
# For each scenario, identify if it's Supervised, Unsupervised, or Reinforcement Learning
scenarios = [
    "Predicting house prices based on size, location, and age",
    "Grouping customers by purchasing behavior without knowing groups beforehand",
    "Teaching a robot to play chess by playing many games",
    "Classifying emails as spam or not spam using labeled examples",
    "Finding hidden topics in news articles without predefined categories"
1
# Your answers (replace 'TYPE' with Supervised, Unsupervised, or Reinforcement)
your_answers = [
    "Supervised",
                      # Scenario 1
    "Unsupervised",
                      # Scenario 2
    "Reinforcement",
                     # Scenario 3
                      # Scenario 4
    "Supervised".
    "Unsupervised"
                      # Scenario 5
# Check answers
correct_answers = ["Supervised", "Unsupervised", "Reinforcement", "Supervised", "Unsupervised"]
print("Scenario Analysis:")
score = 0
for i, (scenario, your_answer, correct) in enumerate(zip(scenarios, your_answers, correct_answers)):
    is_correct = your_answer == correct
    score += is_correct
    status = "✓" if is_correct else "💢"
    print(f"{status} {i+1}. {scenario}")
              Your answer: {your_answer} | Correct: {correct}")
    print()
print(f"Score: {score}/{len(scenarios)} ({score/len(scenarios)*100:.0f}%)")
Assessment Task 1: Identify Machine Learning Types
_____
Scenario Analysis:
1. Predicting house prices based on size, location, and age
   Your answer: Supervised | Correct: Supervised
2. Grouping customers by purchasing behavior without knowing groups beforehand
   Your answer: Unsupervised | Correct: Unsupervised
3. Teaching a robot to play chess by playing many games
   Your answer: Reinforcement | Correct: Reinforcement

✓ 4. Classifying emails as spam or not spam using labeled examples

   Your answer: Supervised | Correct: Supervised
5. Finding hidden topics in news articles without predefined categories
   Your answer: Unsupervised | Correct: Unsupervised
Score: 5/5 (100%)
```

Part 9: Real-World Applications and Case Studies

Let's explore how the concepts we've learned apply to real-world scenarios.

Case Study 1: Recommendation Systems (Netflix, Amazon)

Problem: Suggest movies/products users might like **ML Type:** Hybrid (Supervised + Unsupervised + Reinforcement) **Data:** User ratings, viewing history, product features **Workflow:** Collect data → Build user profiles → Train models → Make recommendations → Learn from feedback

Case Study 2: Fraud Detection (Banks, Credit Cards)

Problem: Identify fraudulent transactions **ML Type**: Supervised Learning (Classification) **Data**: Transaction amounts, locations, times, merchant types **Workflow**: Historical fraud data → Feature engineering → Train classifier → Real-time scoring → Continuous monitoring

Case Study 3: Medical Diagnosis (Healthcare)

Problem: Assist doctors in diagnosing diseases **ML Type:** Supervised Learning (Classification) **Data:** Medical images, patient symptoms, lab results **Workflow:** Labeled medical data \rightarrow Image processing \rightarrow Train deep learning models \rightarrow Clinical validation \rightarrow Deployment with human oversight

Your Turn: Think of Applications

Consider these industries and think about how ML could be applied:

- Transportation: Autonomous vehicles, route optimization
- Agriculture: Crop monitoring, yield prediction
- Education: Personalized learning, automated grading
- · Entertainment: Content creation, game Al

Part 10: Complete ML Workflow Summary

Let's summarize the complete machine learning workflow we've learned:

The Machine Learning Lifecycle

- Problem Definition

 Data Collection & Exploration
 Data Preprocessing & Feature Engineering
 Model Selection & Training
 Model Evaluation & Validation
 Model Deployment & Monitoring
 Continuous Improvement
- Checklist for Every ML Project:

Data Phase:

- Understand the problem and define success metrics
- Collect and explore the dataset
- Check for missing values, outliers, and data quality issues
- Usualize data to understand patterns and relationships

Modeling Phase:

- Split data into training and testing sets
- Select appropriate algorithms for the problem type
- Devaluate using appropriate metrics (accuracy, precision, recall, etc.)

Deployment Phase:

- Validate model performance on new data
- Document the model and its limitations
- Deploy responsibly with monitoring systems
- Plan for model updates and maintenance

6 Key Takeaways:

- 1. Start Simple: Begin with basic models before trying complex ones
- 2. Understand Your Data: EDA is crucial for success
- 3. Validate Properly: Always test on unseen data
- 4. Iterate: ML is an iterative process of improvement
- 5. Document Everything: Keep track of experiments and results

Your Reflection and Analysis

Instructions: Complete the reflection below by editing this markdown cell.

My Understanding of Machine Learning Types

Supervised Learning: Supervised learning is when you train a machine to learn from examples that already have the right answers. **Unsupervised Learning**: Unsupervised learning is when you give the machine data without the answers, and figures out patterns on its own. **Reinforcement Learning**: Reinforcement learning is when a machine learns by trying things out and getting feedback in the form of rewards or penalties

My Analysis of the Wine Classification Project

Best performing model: Decision Tree Classifier

Why do you think this model performed better?: The Decision Tree did better because it can figure out more complicated patterns in the wine data.

What would you try next to improve performance?: I would try adjusting the settings of the Decision Tree to make it more accurate. I'd also test different models and use cross-validation to make sure the results are consistent and not just a lucky split of the data.

Real-World Application Ideas

Industry of Interest: Healthcare

ML Problem: Predicting if a patient might have diabetes based on their medical information.

Type of ML: Supervised Learning

Data Needed: Patient age, weight, blood sugar levels, family history, and other health test results.

Key Learnings

Most important concept learned: I learned how different types of machine learning are used for different kinds of problems.

Most challenging part: The hardest part was understanding why some models perform better than others and how to explain it in simple terms

Questions for further exploration: How do bigger, more advanced models improve on the basic ones?

Lab Summary and Next Steps

- **What You've Accomplished:**
- ✓ Understood ML Types: Supervised, Unsupervised, and Reinforcement Learning
- Mastered ML Workflow: Data → Model → Evaluation → Insights
- **☑ Built Classification Models**: Logistic Regression and Decision Trees
- Evaluated Model Performance: Accuracy, Confusion Matrix, Classification Report
- ✓ Worked with Real Data: Wine dataset analysis and modeling
- ✓ Applied Best Practices: Data splitting, model comparison, interpretation
- Preparation for Module 4:

In the next lab, you'll dive deeper into:

- Exploratory Data Analysis (EDA): Advanced visualization techniques
- · Data Quality Assessment: Handling missing values, outliers, and duplicates
- Statistical Analysis: Understanding distributions and relationships
- Data Storytelling: Communicating insights effectively

Action Items:

- 1. Upload this notebook to your GitHub repository
- 2. Experiment with different features in the wine dataset
- 3. Try other datasets from sklearn.datasets (digits, breast cancer, boston)
- 4. Practice the 6-step ML workflow on a new problem
- 5. Document your experiments and findings

Additional Resources:

- Scikit-learn User Guide
- Machine Learning Mastery
- Kaggle Learn Free micro-courses
- Google's Machine Learning Crash Course

Reflection Questions:

- 1. Which type of machine learning (supervised/unsupervised/reinforcement) interests you most and why?
- 2. What was the most challenging part of the ML workflow for you?
- 3. How might you apply these concepts to a problem in your field of interest?
- 4. What questions do you have about machine learning that you'd like to explore further?

Congratulations on completing Module 3! You've taken a significant step in your machine learning journey. 🏂

Demombers Machine learning is a shill that improves with practice. Vaca asperimenting atom environs and dealt he affect to make mistalian